

Computer Architecture (Practical Class)

Pointers in C

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2025/2026

- A program sees memory as one large array containing bytes
- We do not use the term "index" when referring to a memory location. We use the term **address**
- When you declare a variable, you are reserving one or more **continuous bytes of memory**
- Each declared variable has an *address*, which indicates where the variable data starts in the memory
- In RV32, **addresses are 32 bits long**, ranging from 0 to $2^{32} - 1$ (4 GiB)
- In RV64, addresses are 64 bits long, ranging from 0 to $2^{64} - 1$ (16 EiB)
 - Most current RV64 implementations provide Sv48 virtual-memory support (48-bit VA), while actual physical-address bits are implementation-defined (often 44 bits, i.e. up to 16 TiB of RAM, or 52 bits, i.e. up to 4 PiB)

- C has special variables, called **pointers**, that are used to store memory addresses
- Pointers are declared like normal variables, with a type associated to it (we will see how this is used later in this class)

Pointers always have the same size

- The size of an address of the underlying architecture (32 bits (4 bytes) in RV32)
- Pointers allow direct access to memory, making it possible to change the values in the memory addresses stored in the pointers
- Some tasks are easier/more efficient to implement when using pointers, and some (such as dynamic memory allocation) are only possible using pointers

- Just like any variable in C, a pointer must be declared before being used
- A pointer is declared using type `*` before the variable identifier:

```
int *ptr1;    // declares a pointer to an integer
char *ptr2;   // declares a pointer to a char
```

- To obtain the address of a variable, use `'&'` before its identifier:

```
int x;
char c;
ptr1 = &x; // store the address of x in ptr1
ptr2 = &c; // store the address of c in ptr2
```

- The *dereference operator* `'*'` accesses the value at a memory address:


```
*ptr1 = 10; // assign 10 to the value pointed by ptr1
*ptr2 = 'X'; // assign 'X' to the value pointed by ptr2
```

Content of a variable and its address (1/2)


Important note

For the sake of simplicity, in the following schemes, an address is represented in only one byte. DO NOT forget that real addresses **always occupy 4 bytes in RV32**

```
char sum;  
char *ptr;  
sum = 10;
```




```
ptr = &sum;
```




Content of a variable and its address (2/2)

- To access the contents of a memory address in a pointer, use the * operator:



```
*ptr = 20;
```



Pointers should always be initialized to a valid address before being used

- Using an uninitialized pointer has undefined behaviour!

- Wrong way¹:

```
int x;  
int *ptr;    // uninitialized pointer  
  
*ptr = 22;    // this is a very bad idea!
```

- Correct way:

```
int x;  
int *ptr;  
  
ptr = &x;    // now is initialized to the address of x  
*ptr = 22;    // assign 22 to the memory pointed by ptr (x)
```

¹This code will, most likely, result in a segmentation fault. A segmentation fault occurs when a program tries to access an invalid memory address. The operating system detects this and terminates the program.

- We have said that the memory is a large array of bytes... So, how do we store data larger than 1 byte?
- Easy: we divide the data into bytes and store it! But, this means we have **two ways of storing data in memory**
- Big Endian
 - Store the most significant byte in the smallest address.
 - Adopted in platforms by Sun, PowerPC Mac, transferring data on the Internet, ...
- Little Endian
 - Store the least significant byte in the smallest address.
 - Adopted by x86, ARM, RISC-V, ...

Little Endian example

- Let us assume we want to store the number 305419896 ($0x12345678$):

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0x12								0x34								0x56								0x78							

- Byte by byte

7	6	5	4	3	2	1	0		
0x78								}	Address x
0x56									Address $x + 1$
0x34									Address $x + 2$
0x12									Address $x + 3$

Big Endian example

- Let us assume we want to store the number 305419896 ($0x12345678$):

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0x12								0x34								0x56								0x78							

- Byte by byte

7	6	5	4	3	2	1	0	
0x12								} Address x } Address $x + 1$ } Address $x + 2$ } Address $x + 3$
0x34								
0x56								
0x78								

Data representation in RV32

```
#include <stdio.h>


int main(){
    unsigned int i = 0xFFAAEEBB;
    short b      = 0x1234;
    char c       = 'A';
    unsigned int x = 0x12345678;

    // %p: Pointer address format specifier
    printf("i em %p\n",&i);
    printf("b em %p\n",&b);
    printf("c em %p\n",&c);
    printf("x em %p\n",&x);
    return 0;
}
```


Variable	Address	Memory (byte by byte in hexadecimal)
i	0x9854fa7c	bb ee aa ff
b	0x9854fa7a	34 12
c	0x9854fa79	41
x	0x9854fa74	78 56 34 12

Handling arrays with pointers


```
char vec[4] = {1,2,3,4};  
char *ptr = NULL;
```



```
ptr = vec; /* Array names already represent the address of their  
           first element, so using '&' is unnecessary */
```



Pointer arithmetic (1/2)



- You can do **pointer arithmetic**:

```
ptr++;           // ptr now is a pointer to the second element of vec
printf("%hhd", *ptr); // will print "2"

ptr+=2;         // ptr now is a pointer to the fourth element of vec
printf("%hhd", *ptr); // will print "4"
```

Important note

The type of the pointer is relevant when performing pointer arithmetic (the compiler decides on the number of bytes added or subtracted based on the pointer type)

```
int *total_ptr;           // total_ptr++ moves the pointer by sizeof(int)
                          // (4 bytes)


char *name_ptr;           // name_ptr++ moves the pointer by sizeof(char)
                          // (1 byte)

short *count_ptr;         // count_ptr++ moves the pointer by sizeof(short)
                          // (2 bytes)


long long *big_ptr;       // big_ptr++ moves the pointer by sizeof(long long)
                          // (8 bytes)
```

Pointer arithmetic: Example (1/2)

```
int vec[4] = {1,2,3,4};  
int *ptr = vec;
```




```
ptr = ptr+2;
```




Pointer arithmetic: Example (2/2)


```
int vec[4] = {1,2,3,4};  
char *ptr = (char *)vec; // a cast is needed to avoid a warning...
```



```
ptr = ptr+1;
```



```
*ptr = 'A';
```



- Active learning activity: What is the output of the following code ?
 - A. 1 2 3 4 5
 - B. 2 3 4 5 6
 - C. 1 3 5 7 9
 - D. None of the above.

Listing 1: pointer-arith.c

```
#include <stdio.h>

int main(){
    char arr[] = {1,2,3,4,5,6,7,8,9,10};
    short *ptr = (short*)arr; // cast is needed to avoid a warning...
    int i;

    for (i=0; i<5; i++) {
        arr[i] = *( ptr + i );
        printf("%hhd ", arr[i]);
    }
    return 0;
}
```

Listing 2: output.c

```
#include <stdio.h>

int main(){
    int a = 2, b = 3;
    int *c = &a, *d = &b;

    printf("Value: %d\n", *c);                // prints the value 2

    printf("Address: %p (%p)\n", &a, c);      // prints the address of a

    if (*c == *d) puts("Same value");        // false

    *c = 3;
    if (*c == *d) puts("Now same value");    // true

    c = d;
    if (c == d) puts ("Now same address");  // true

    return 0;
}
```

Listing 3: input.c

```

#include <stdio.h>
#include <stdlib.h>

int main(){
    int n1 = -1, n2 = -1, n3 = -1;
    int *ptr = &n2;
    char buf[BUFSIZ], c;

    printf ("Enter a number: ");
    scanf("%d", &n1);
    printf ("Enter another number: ");
    scanf("%[^0-9]%d", buf, ptr);           // garbage goes to buf

    /* You will notice that scanf is a bit problematic...
       An option, is to read a string and convert as needed
       First, flush characters scanf() leaves in the input */
    while((c = getchar()) != '\n' && c != EOF);
    printf ("Enter yet another number: ");
    if (fgets(buf, sizeof(buf), stdin) != NULL) {
        n3 = atoi(buf);                     // returns 0 if conversion fails
        printf ("You entered %d %d %d\n", n1, n2, n3);
        return EXIT_SUCCESS;
    }
    return EXIT_FAILURE;                   // if reaches here, fgets() failed
}

```

- Write the representation in memory, in Big endian and Little endian of the following values:
 - 0x1188 (16 bits)
 - 0xff3455b6 (32 bits)
 - 0x28934def (32 bits)
- Correct the following code:

Listing 4: exerc01.c

```
int main () {  
    int * ptr ;  
    int i;  
    int sum=0;  
  
    for(i=0; i<10; i++){  
        scanf("%d", ptr);  
        sum = sum + *ptr;  
    }  
  
    printf("Sum = %d \n", sum);  
  
    return 0;  
}
```